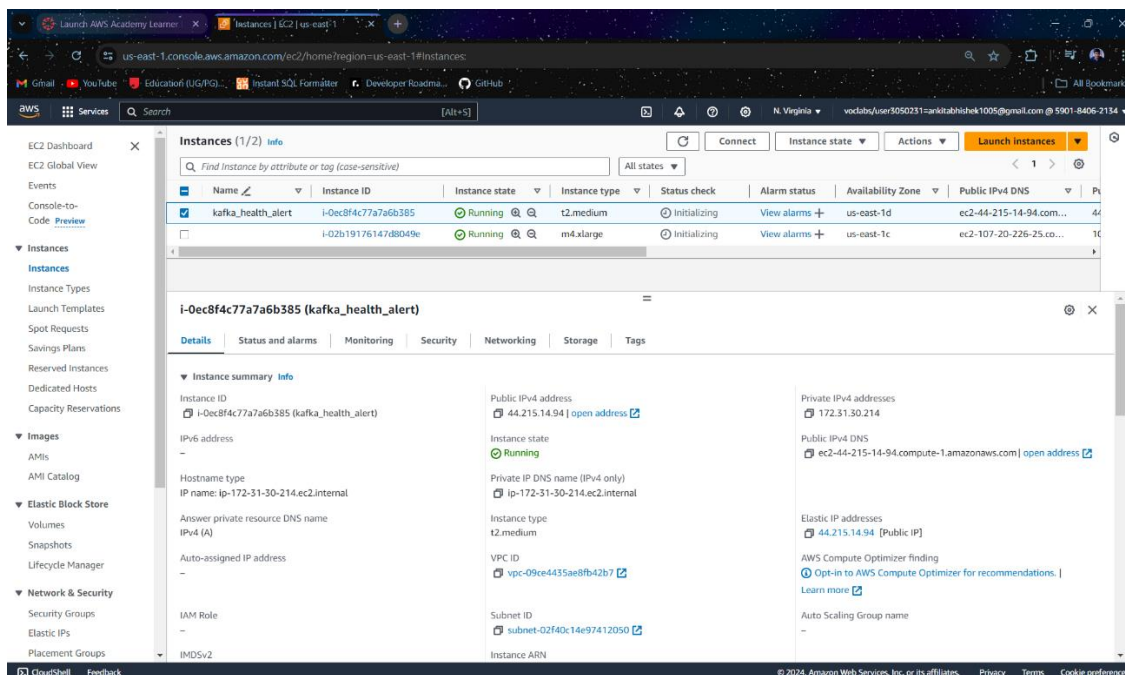
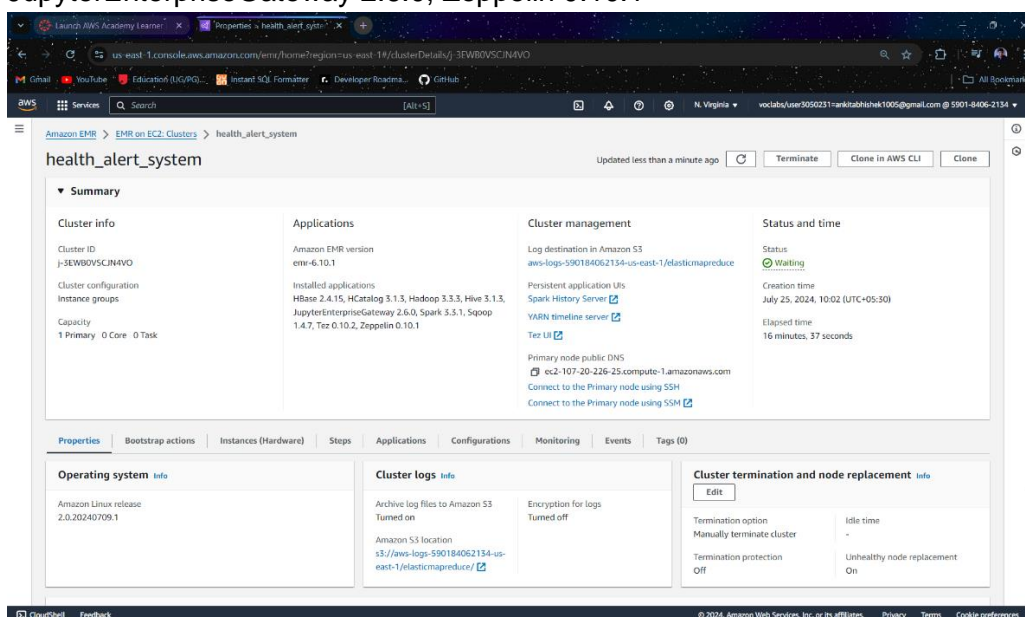


## CREATION OF KAFKA CLUSTER AND EMR CLUSTER

1. Create a kafka cluster with the help of ec2 instance. (Referred with PDF attached in the Apache Kafka modules). Make the required configuration needed to run kafka instance. Kafka is pre-installed on below mentioned ec2 cluster with the selection of **ami-06c41d8b5a6ddd3c2** while creating **Amazon Machine Image** as pdf within modules.



2. Create an EMR instance with required below mentioned libraries (Referred with PDF attached in the modules).  
Spark 3.3.1, Sqoop 1.4.7, HBase 2.4.15, HCatalog 3.1.3, Hadoop 3.3.3, Hive 3.1.3, JupyterEnterpriseGateway 2.6.0, Zeppelin 0.10.1



## INSTALLING REQUIRED PACKAGES ON KAFKA CLUSTER

Sudo pip3 install kafka-python  
Sudo pip3 install mysql-connector  
Sudo pip3 install boto3

```
ec2-user@ip-172-31-30-214:~  
[ec2-user@ip-172-31-30-214 ~]$ sudo pip3 install kafka-python  
WARNING: Running pip install with root privileges is generally not a good idea.  
Try `pip3 install --user` instead.  
Collecting kafka-python  
  Downloading kafka-python-2.0.2-py2.py3-none-any.whl (246 kB)  
    |████████████████████| 246 kB 36.2 MB/s  
Installing collected packages: kafka-python  
Successfully installed kafka-python-2.0.2  
[ec2-user@ip-172-31-30-214 ~]$ sudo pip3 install mysql-connector  
WARNING: Running pip install with root privileges is generally not a good idea.  
Try `pip3 install --user` instead.  
Collecting mysql-connector  
  Downloading mysql-connector-2.2.9.tar.gz (11.9 MB)  
    |████████████████████| 11.9 MB 69 kB/s  
Using legacy 'setup.py install' for mysql-connector, since package 'wheel' is not installed.  
Installing collected packages: mysql-connector  
  Running setup.py install for mysql-connector ... done  
Successfully installed mysql-connector-2.2.9  
[ec2-user@ip-172-31-30-214 ~]$ sudo pip3 install boto3  
WARNING: Running pip install with root privileges is generally not a good idea.  
Try `pip3 install --user` instead.  
Collecting boto3  
  Downloading boto3-1.33.13-py3-none-any.whl (139 kB)  
    |████████████████████| 139 kB 13.2 MB/s  
Try `pip3 install --user` instead.  
Collecting boto3  
  Downloading boto3-1.33.13-py3-none-any.whl (139 kB)  
    |████████████████████| 139 kB 13.2 MB/s  
Collecting s3transfer<0.9.0,>=0.8.2  
  Downloading s3transfer-0.8.2-py3-none-any.whl (82 kB)  
    |████████████████████| 82 kB 122 kB/s  
Collecting jmespath<2.0.0,>=0.7.1  
  Downloading jmespath-1.0.1-py3-none-any.whl (20 kB)  
Collecting botocore<1.34.0,>=1.33.13  
  Downloading botocore-1.33.13-py3-none-any.whl (11.8 MB)  
    |████████████████████| 11.8 MB 35 kB/s  
Collecting python-dateutil<3.0.0,>=2.1  
  Downloading python-dateutil-2.9.0.post0-py2.py3-none-any.whl (229 kB)  
    |████████████████████| 229 kB 59.4 MB/s  
Collecting urllib3<1.27,>=1.25.4; python_version < "3.10"  
  Downloading urllib3-1.26.19-py2.py3-none-any.whl (143 kB)  
    |████████████████████| 143 kB 63.5 MB/s  
Collecting six>=1.5  
  Downloading six-1.16.0-py2.py3-none-any.whl (11 kB)  
Installing collected packages: six, python-dateutil, urllib3, jmespath, botocore, s3transfer, boto3  
Successfully installed boto3-1.33.13 botocore-1.33.13 jmespath-1.0.1 python-dateutil-2.9.0.post0 s3transfer-0.8.2 six-1.16.0 urllib3-1.26.19  
[ec2-user@ip-172-31-30-214 ~]$
```

## STATEMENT FOR STARTING KAFKA SERVER

### 1. STARTING ZOOKEEPER SERVER:

Inside cd downloads/kafka\_2.12-2.3.0 run

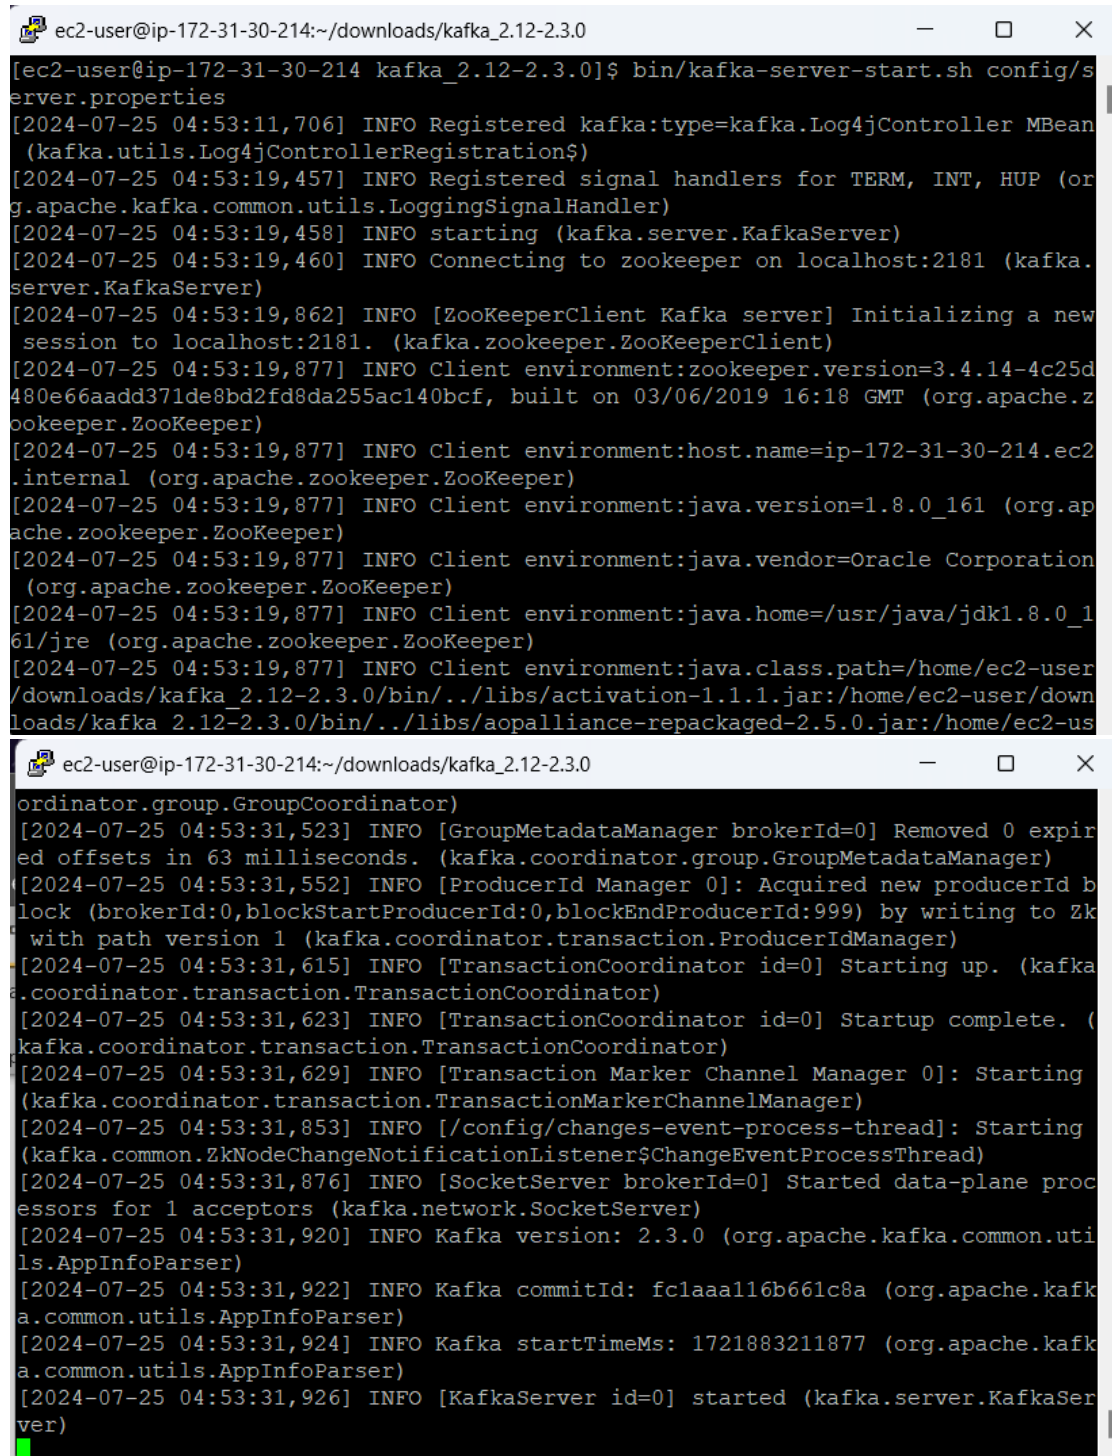
**bin/zookeeper-server-start.sh config/zookeeper.properties**

```
ec2-user@ip-172-31-30-214:~/downloads/kafka_2.12-2.3.0
[ec2-user@ip-172-31-30-214 kafka_2.12-2.3.0]$ bin/zookeeper-server-start.sh config/zookeeper.properties
[2024-07-25 04:53:01,268] INFO Reading configuration from: config/zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2024-07-25 04:53:01,457] INFO autopurge.snapRetainCount set to 3 (org.apache.zookeeper.server.DataDirCleanupManager)
[2024-07-25 04:53:01,457] INFO autopurge.purgeInterval set to 0 (org.apache.zookeeper.server.DataDirCleanupManager)
[2024-07-25 04:53:01,458] INFO Purge task is not scheduled. (org.apache.zookeeper.server.DataDirCleanupManager)
[2024-07-25 04:53:01,458] WARN Either no config or no quorum defined in config, running in standalone mode (org.apache.zookeeper.server.quorum.QuorumPeerMain)
[2024-07-25 04:53:01,560] INFO Reading configuration from: config/zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2024-07-25 04:53:01,560] INFO Starting server (org.apache.zookeeper.server.ZooKeeperServerMain)
[2024-07-25 04:53:01,683] INFO Server environment:zookeeper.version=3.4.14-4c25d480e66aadd371de8bd2fd8da255ac140bcf, built on 03/06/2019 16:18 GMT (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:01,683] INFO Server environment:host.name=ip-172-31-30-214.ec2.internal (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:01,683] INFO Server environment:java.version=1.8.0_161 (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:01,684] INFO Server environment:java.vendor=Oracle Corporation (org.apache.zookeeper.server.ZooKeeperServer)
```

```
ec2-user@ip-172-31-30-214:~/downloads/kafka_2.12-2.3.0
[2024-07-25 04:53:01,685] INFO Server environment:java.compiler=<NA> (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:01,685] INFO Server environment:os.name=Linux (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:01,685] INFO Server environment:os.arch=amd64 (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:01,685] INFO Server environment:os.version=4.14.193-149.317.amzn2.x86_64 (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:01,685] INFO Server environment:user.name=ec2-user (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:01,685] INFO Server environment:user.home=/home/ec2-user (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:01,685] INFO Server environment:user.dir=/home/ec2-user/downloads/kafka_2.12-2.3.0 (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:02,135] INFO tickTime set to 3000 (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:02,135] INFO minSessionTimeout set to -1 (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:02,136] INFO maxSessionTimeout set to -1 (org.apache.zookeeper.server.ZooKeeperServer)
[2024-07-25 04:53:02,760] INFO Using org.apache.zookeeper.server.NIOServerCnxnFactory as server connection factory (org.apache.zookeeper.server.ServerCnxnFactory)
[2024-07-25 04:53:03,150] INFO binding to port 0.0.0.0/0.0.0.0:2181 (org.apache.zookeeper.server.NIOServerCnxnFactory)
```

## 2. STARTING KAFKA SERVER:

Into another putty Session of kafka cluster inside cd downloads/kafka\_2.12-2.3.0 run **bin/kafka-server-start.sh config/server.properties**



```
ec2-user@ip-172-31-30-214:~/downloads/kafka_2.12-2.3.0
[ec2-user@ip-172-31-30-214 kafka_2.12-2.3.0]$ bin/kafka-server-start.sh config/s
server.properties
[2024-07-25 04:53:11,706] INFO Registered kafka:type=kafka.Log4jController MBean
(kafka.utils.Log4jControllerRegistration$)
[2024-07-25 04:53:19,457] INFO Registered signal handlers for TERM, INT, HUP (or
g.apache.kafka.common.utils.LoggingSignalHandler)
[2024-07-25 04:53:19,458] INFO starting (kafka.server.KafkaServer)
[2024-07-25 04:53:19,460] INFO Connecting to zookeeper on localhost:2181 (kafka.
server.KafkaServer)
[2024-07-25 04:53:19,862] INFO [ZooKeeperClient Kafka server] Initializing a new
session to localhost:2181. (kafka.zookeeper.ZooKeeperClient)
[2024-07-25 04:53:19,877] INFO Client environment:zookeeper.version=3.4.14-4c25d
480e66aadd371de8bd2fd8da255ac140bcf, built on 03/06/2019 16:18 GMT (org.apache.z
ookeeper.ZooKeeper)
[2024-07-25 04:53:19,877] INFO Client environment:host.name=ip-172-31-30-214.ec2
.internal (org.apache.zookeeper.ZooKeeper)
[2024-07-25 04:53:19,877] INFO Client environment:java.version=1.8.0_161 (org.ap
ache.zookeeper.ZooKeeper)
[2024-07-25 04:53:19,877] INFO Client environment:java.vendor=Oracle Corporation
(org.apache.zookeeper.ZooKeeper)
[2024-07-25 04:53:19,877] INFO Client environment:java.home=/usr/java/jdk1.8.0_1
61/jre (org.apache.zookeeper.ZooKeeper)
[2024-07-25 04:53:19,877] INFO Client environment:java.class.path=/home/ec2-user
/downloads/kafka_2.12-2.3.0/bin/../libs/activation-1.1.1.jar:/home/ec2-user/down
loads/kafka_2.12-2.3.0/bin/../libs/aopalliance-repackaged-2.5.0.jar:/home/ec2-us
ordinator.group.GroupCoordinator)
[2024-07-25 04:53:31,523] INFO [GroupMetadataManager brokerId=0] Removed 0 expir
ed offsets in 63 milliseconds. (kafka.coordinator.group.GroupMetadataManager)
[2024-07-25 04:53:31,552] INFO [ProducerId Manager 0]: Acquired new producerId b
lock (brokerId=0,blockStartProducerId=0,blockEndProducerId=999) by writing to Zk
with path version 1 (kafka.coordinator.transaction.ProducerIdManager)
[2024-07-25 04:53:31,615] INFO [TransactionCoordinator id=0] Starting up. (kafka
.coordinator.transaction.TransactionCoordinator)
[2024-07-25 04:53:31,623] INFO [TransactionCoordinator id=0] Startup complete. (
kafka.coordinator.transaction.TransactionCoordinator)
[2024-07-25 04:53:31,629] INFO [Transaction Marker Channel Manager 0]: Starting
(kafka.coordinator.transaction.TransactionMarkerChannelManager)
[2024-07-25 04:53:31,853] INFO [/config/changes-event-process-thread]: Starting
(kafka.common.ZkNodeChangeNotificationListener$ChangeEventProcessThread)
[2024-07-25 04:53:31,876] INFO [SocketServer brokerId=0] Started data-plane proc
essors for 1 acceptors (kafka.network.SocketServer)
[2024-07-25 04:53:31,920] INFO Kafka version: 2.3.0 (org.apache.kafka.common.uti
ls.AppInfoParser)
[2024-07-25 04:53:31,922] INFO Kafka commitId: fclaa116b661c8a (org.apache.kafk
a.common.utils.AppInfoParser)
[2024-07-25 04:53:31,924] INFO Kafka startTimeMs: 1721883211877 (org.apache.kafk
a.common.utils.AppInfoParser)
[2024-07-25 04:53:31,926] INFO [KafkaServer id=0] started (kafka.server.KafkaSer
ver)
```



## STATEMENT TO CREATE TOPICS

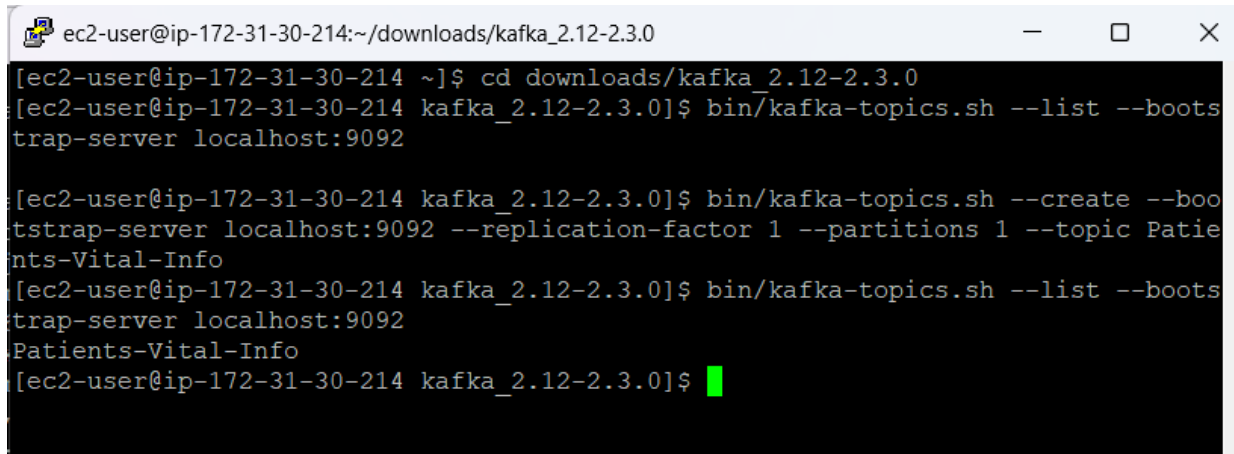
To create topic in kafka server, the command used is

```
bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1  
--partitions 1 --topic Patients-Vital-Info
```

## STATEMENT TO LIST TOPICS

To list the created topic inside cd downloads/kafka\_2.12-2.3.0, the command used is

```
bin/kafka-topics.sh --list --bootstrap-server localhost:9092
```



```
ec2-user@ip-172-31-30-214:~/downloads/kafka_2.12-2.3.0  
[ec2-user@ip-172-31-30-214 ~]$ cd downloads/kafka_2.12-2.3.0  
[ec2-user@ip-172-31-30-214 kafka_2.12-2.3.0]$ bin/kafka-topics.sh --list --bootstrap-server localhost:9092  
  
[ec2-user@ip-172-31-30-214 kafka_2.12-2.3.0]$ bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --topic Patients-Vital-Info  
[ec2-user@ip-172-31-30-214 kafka_2.12-2.3.0]$ bin/kafka-topics.sh --list --bootstrap-server localhost:9092  
Patients-Vital-Info  
[ec2-user@ip-172-31-30-214 kafka_2.12-2.3.0]$
```

## EXECUTING PRODUCER APPLICATION AND CONSUMER APPLICATION:

Producer application which is file named as **kafka\_produce\_patient\_vitals.py** is built on the **python language** which will consume data residing on rds with below mentioned credentials:

```
Hostname = "upgradtest.cyaieic9bmnf.us-east-1.rds.amazonaws.com"
```

```
username = "student"
```

```
password = "STUDENT123"
```

```
dbname = "testdatabase".
```

Consumer Application which is file named as **kafka\_spark\_patient\_vitals.py** is built on the **Apache PySpark** language which will consume data being produced with the help of above mentioned producer application

**NOTE:** Run the producer application on ec2 Kafka cluster after starting the consumer application on EMR cluster created with Spark, Hive and another libraries

## STATEMENT FOR EXECUTING PRODUCER APPLICATION AND CONSUMER APPLICATION

### Spark Submitting Job to Consume Message from The Topic Patients-Vital-Info And Stored To HDFS Location

For Producer application: **python3 kafka\_produce\_patients\_vitals.py**

```

[ec2-user@ip-172-31-30-214:~]$ ls
anaconda2          jdk-8u161-linux-x64.tar.gz
Anaconda2-4.1.1-Linux-x86_64.sh  kafka_produce_patient_vitals.py
downloads          Notebook
[ec2-user@ip-172-31-30-214:~]$ python3 kafka_produce_patient_vitals.py

[hadoop@ip-172-31-90-234:~]$ ls
kafka_spark_patient_vitals.py
[hadoop@ip-172-31-90-234:~]$ spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.3.1 kafka_spark_patient_vitals.py

[ec2-user@ip-172-31-30-214:~]$ ls
anaconda2          jdk-8u161-linux-x64.tar.gz
Anaconda2-4.1.1-Linux-x86_64.sh  kafka_produce_patient_vitals.py
downloads          Notebook
[ec2-user@ip-172-31-30-214:~]$ python3 kafka_produce_patient_vitals.py
{'customerId': 1, 'heartBeat': 74, 'bp': 202}
{'customerId': 2, 'heartBeat': 68, 'bp': 173}
{'customerId': 3, 'heartBeat': 71, 'bp': 152}

-234.ec2.internal:46147
24/07/25 05:16:41 INFO BlockManager: Using org.apache.spark.storage.RandomBlockReplicationPolicy for block replication policy
24/07/25 05:16:41 INFO BlockManager: external shuffle service port = 7337
24/07/25 05:16:41 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-31-90-234.ec2.internal, 46147, None)
24/07/25 05:16:41 INFO BlockManagerMasterEndpoint: Registering block manager ip-172-31-90-234.ec2.internal:46147 with 912.3 MiB RAM, BlockManagerId(driver, ip-172-31-90-234.ec2.internal, 46147, None)
24/07/25 05:16:41 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-31-90-234.ec2.internal, 46147, None)
24/07/25 05:16:41 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, ip-172-31-90-234.ec2.internal, 46147, None)
24/07/25 05:16:42 INFO SingleEventLogFileWriter: Logging events to hdfs://var/log/spark/apps/local-1721884601542.inprogress
-----
Batch: 0
-----
+-----+-----+-----+
|customerId|heartBeat|bp|message_time|
+-----+-----+-----+

```

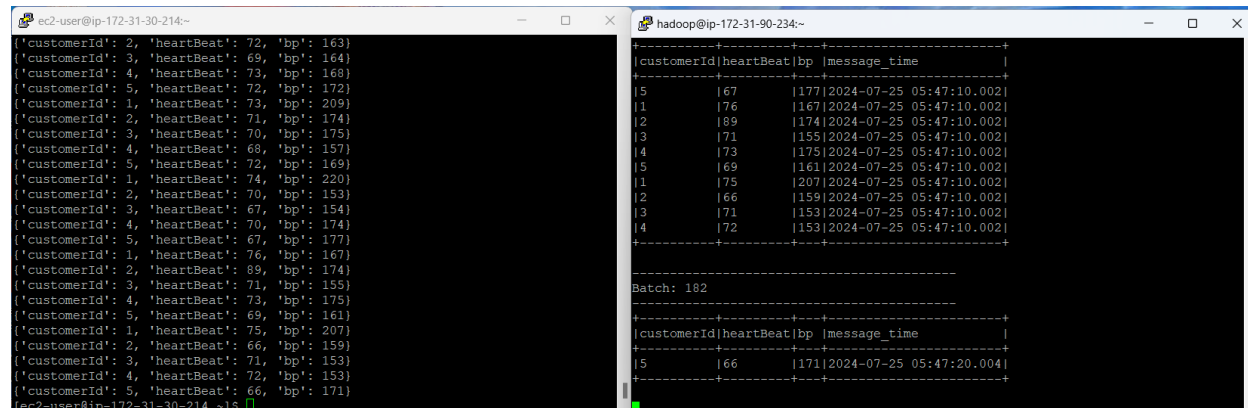
For Consumer Application: **spark-submit --packages org.apache.spark:spark-sql-kafka-0-10\_2.12:3.3.1 kafka\_spark\_patient\_vitals.py**

```
hadoop@ip-172-31-90-129:~
EE:::::EEEEEEEE::::E M:::::M      M:::::M   R:::R      R:::::R
E:::::E M:::::M      M:::::M   RR:::R      R:::::R
EEEEEEEEEEEEEEEEEEEE MMMMMMM   MMMMMMM   RRRRRRR   RRRRRR

[hadoop@ip-172-31-90-129 ~]$ ls
[hadoop@ip-172-31-90-129 ~]$ ls
kafka_spark_patient_vitals.py
[hadoop@ip-172-31-90-129 ~]$ spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.3.1 kafka_spark_patient_vitals.py
:: loading settings :: url = jar:file:/usr/lib/spark/jars/ivy-2.5.0.jar!/org/apache/ivy/core/settings/ivysettings.xml
Ivy Default Cache set to: /home/hadoop/.ivy2/cache
The jars for the packages stored in: /home/hadoop/.ivy2/jars
org.apache.spark#spark-sql-kafka-0-10_2.12 added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-parent-3252d621-c6ec-4b42-ba7b-10c52d78769d;1.0
   confs: [default]
   found org.apache.spark#spark-sql-kafka-0-10_2.12;3.3.1 in central
   found org.apache.spark#spark-token-provider-kafka-0-10_2.12;3.3.1 in central
   found org.apache.kafka#kafka-clients;2.8.1 in central
   found org.lz4#lz4-java;1.8.0 in central
   found org.xerial.snappy#snappy-java;1.1.8.4 in central
```

```
hadoop@ip-172-31-90-129:~
-129.ec2.internal:45885
24/07/21 10:19:50 INFO BlockManager: Using org.apache.spark.storage.RandomBlockReplicationPolicy for block replication policy
24/07/21 10:19:50 INFO BlockManager: external shuffle service port = 7337
24/07/21 10:19:50 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-31-90-129.ec2.internal, 45885, None)
24/07/21 10:19:50 INFO BlockManagerMasterEndpoint: Registering block manager ip-172-31-90-129.ec2.internal:45885 with 912.3 MiB RAM, BlockManagerId(driver, ip-172-31-90-129.ec2.internal, 45885, None)
24/07/21 10:19:50 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-31-90-129.ec2.internal, 45885, None)
24/07/21 10:19:50 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, ip-172-31-90-129.ec2.internal, 45885, None)
24/07/21 10:19:51 INFO SingleEventLog FileWriter: Logging events to hdfs:/var/log/spark/apps/local-1721557189930.inprogress
-----
Batch: 0
-----
+-----+-----+---+-----+
|customerId|heartBeat|bp |message_time|
+-----+-----+---+-----+
+-----+-----+---+-----+
```

After 30 minutes when all 1800 data being streamed and saved to Parquet file of the required HDFS location



The image shows two terminal windows. The left window, titled 'ec2-user@ip-172-31-30-214:~', displays a continuous stream of JSON data records. The right window, titled 'hadoop@ip-172-31-90-234:~', shows a table of data with columns 'customerId', 'heartBeat', 'bp', and 'message\_time'. Below the table, it indicates 'Batch: 182' and shows a single record from that batch.

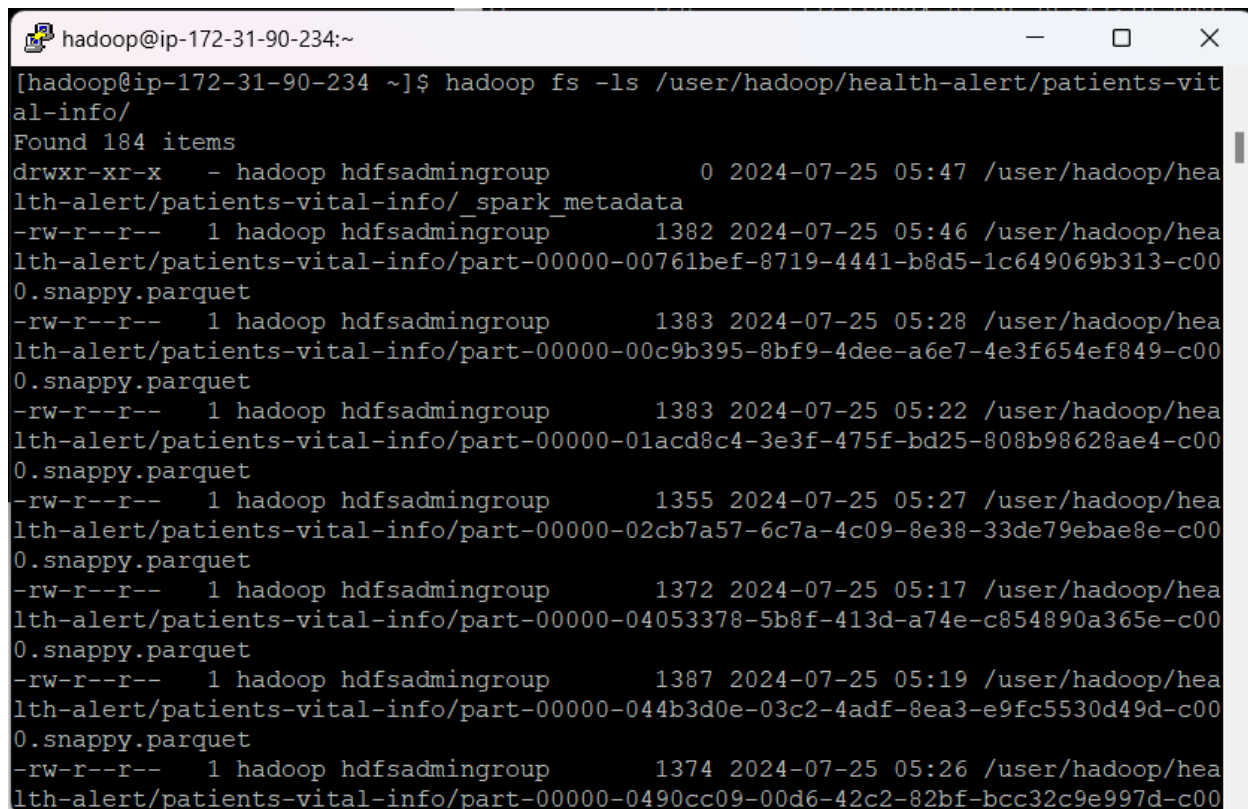
```

ec2-user@ip-172-31-30-214:~$ cat /dev/stdin
{"customerId": 2, "heartBeat": 72, "bp": 163}
{"customerId": 3, "heartBeat": 69, "bp": 164}
{"customerId": 4, "heartBeat": 73, "bp": 168}
{"customerId": 5, "heartBeat": 72, "bp": 172}
{"customerId": 1, "heartBeat": 73, "bp": 209}
{"customerId": 2, "heartBeat": 71, "bp": 174}
{"customerId": 3, "heartBeat": 70, "bp": 175}
{"customerId": 4, "heartBeat": 68, "bp": 157}
{"customerId": 5, "heartBeat": 72, "bp": 169}
{"customerId": 1, "heartBeat": 74, "bp": 220}
{"customerId": 2, "heartBeat": 70, "bp": 153}
{"customerId": 3, "heartBeat": 67, "bp": 154}
{"customerId": 4, "heartBeat": 70, "bp": 174}
{"customerId": 5, "heartBeat": 67, "bp": 177}
{"customerId": 1, "heartBeat": 76, "bp": 167}
{"customerId": 2, "heartBeat": 89, "bp": 174}
{"customerId": 3, "heartBeat": 71, "bp": 155}
{"customerId": 4, "heartBeat": 73, "bp": 175}
{"customerId": 5, "heartBeat": 69, "bp": 161}
{"customerId": 1, "heartBeat": 75, "bp": 207}
{"customerId": 2, "heartBeat": 66, "bp": 159}
{"customerId": 3, "heartBeat": 71, "bp": 153}
{"customerId": 4, "heartBeat": 72, "bp": 153}
{"customerId": 5, "heartBeat": 66, "bp": 171}

hadoop@ip-172-31-90-234:~$ cat /dev/stdin
customerId|heartBeat|bp|message_time|
+-----+
5|67|177|2024-07-25 05:47:10.002|
1|76|167|2024-07-25 05:47:10.002|
2|89|174|2024-07-25 05:47:10.002|
3|71|155|2024-07-25 05:47:10.002|
4|73|175|2024-07-25 05:47:10.002|
5|69|161|2024-07-25 05:47:10.002|
1|75|207|2024-07-25 05:47:10.002|
2|66|159|2024-07-25 05:47:10.002|
3|71|153|2024-07-25 05:47:10.002|
4|72|153|2024-07-25 05:47:10.002|
+-----+
Batch: 182
+-----+
customerId|heartBeat|bp|message_time|
+-----+
5|66|171|2024-07-25 05:47:20.004|
+-----+
  
```

## STATEMENT TO CHECK DATA STORED IN HDFS LOCATION:

**hadoop fs -ls /user/hadoop/health-alert/patients-vital-info/**



The image shows a terminal window titled 'hadoop@ip-172-31-90-234:~' with the command 'hadoop fs -ls /user/hadoop/health-alert/patients-vital-info/' executed. The output shows 184 items, including a directory for '\_spark\_metadata' and several Parquet files.

```

hadoop@ip-172-31-90-234:~$ hadoop fs -ls /user/hadoop/health-alert/patients-vital-info/
Found 184 items
drwxr-xr-x  - hadoop hdfsadmin 0 2024-07-25 05:47 /user/hadoop/health-alert/patients-vital-info/_spark_metadata
-rw-r--r--  1 hadoop hdfsadmin 1382 2024-07-25 05:46 /user/hadoop/health-alert/patients-vital-info/part-00000-00761bef-8719-4441-b8d5-1c649069b313-c000.snappy.parquet
-rw-r--r--  1 hadoop hdfsadmin 1383 2024-07-25 05:28 /user/hadoop/health-alert/patients-vital-info/part-00000-00c9b395-8bf9-4dee-a6e7-4e3f654ef849-c000.snappy.parquet
-rw-r--r--  1 hadoop hdfsadmin 1383 2024-07-25 05:22 /user/hadoop/health-alert/patients-vital-info/part-00000-01acd8c4-3e3f-475f-bd25-808b98628ae4-c000.snappy.parquet
-rw-r--r--  1 hadoop hdfsadmin 1355 2024-07-25 05:27 /user/hadoop/health-alert/patients-vital-info/part-00000-02cb7a57-6c7a-4c09-8e38-33de79ebae8e-c000.snappy.parquet
-rw-r--r--  1 hadoop hdfsadmin 1372 2024-07-25 05:17 /user/hadoop/health-alert/patients-vital-info/part-00000-04053378-5b8f-413d-a74e-c854890a365e-c000.snappy.parquet
-rw-r--r--  1 hadoop hdfsadmin 1387 2024-07-25 05:19 /user/hadoop/health-alert/patients-vital-info/part-00000-044b3d0e-03c2-4adf-8ea3-e9fc5530d49d-c000.snappy.parquet
-rw-r--r--  1 hadoop hdfsadmin 1374 2024-07-25 05:26 /user/hadoop/health-alert/patients-vital-info/part-00000-0490cc09-00d6-42c2-82bf-bcc32c9e997d-c000.snappy.parquet
  
```



## STATEMENT TO READ ONE OF THE FILES USING '-CAT'

```
hadoop fs -cat /user/hadoop/health-alert/patients-vital-info/part-00000-ffcd45dc-ef2e-4219-b11e-f26f0e051d73-c000.snappy.parquet
```

```
hadoop@ip-172-31-90-234:~$ hadoop fs -cat /user/hadoop/health-alert/patients-vital-info/part-00000-feb44dbe-51f1-44f3-b5ae-c023c3f6104e-c000.snappy.parquet
```

```
PAR1(,  
L FD#(  
LJHBGE  
4l  
\X
```

```
<  
BJZN6\H  
spark_schem%  
customerId% heartBeat%bp%  
message_timeL&25  
customerId&Z(,.,&5 heartBeat&&JB(JB,.,&bp&<(.,&%  
message_timepx&&  
@7Ne%  
@7Ne%(  
@7Ne%  
@7Ne%,org.apache.spark.version3.3.1)org.apache.spark.sql.parquet.row.metad  
ata={"type":"struct","fields":[{"name":"customerId","type":"integer","nullable":t  
rue,"metadata":{"name":"bp","type":"integer","nullable":true,"metadat  
a":{"name":"time","type":"timestamp","nullable":false,"metadata":{"name":"version  
1.12.2 (build 88690eb334b5f0273c2b37d8d767559f594bf245)L PAR1PuTTYPuTTYPuTTYPuTTY  
PuTTYPuTTYPuTTYPuTTYPuTTYPuTTYPuTTYPuTTY[hadoop@ip-172-31-90-234 ~]$ PuTTYPuTTYPuTTYPu
```